

GenomeX Assurance Platform

Product Whitepaper

How the product works, what buyers get, and what evidence it produces.

Product paper

Buyer-facing solution whitepaper for the GenomeX Assurance Platform.

Why Enterprise Assurance exists is the subject of When Data Becomes Infrastructure. The architecture this paper inherits is specified in GenomeX Assurance Platform.

Version	3.0 — September 2026
Audience	CISO, CRO, backup and DR leaders, AI platform owners, compliance officers, MSSPs, cyber insurers, investors, OEM partners
Scope	Product surfaces, domain gates, receipts, deployment, and maturity
Architecture owner	GenomeX Assurance Platform (analyst / architecture paper)

Availability legend

Badge	Meaning
Available	Shipped or in active production or lab validation
Pilot	Implemented; partner or production sign-off in progress
Roadmap	Specified; not yet generally available

Executive summary

Enterprises run on five classes of digital assets that must be trusted before use: **recovery chains** (environment plus recovery assets), **knowledge** (corpora feeding RAG and search), **training data**, **software releases**, and **autonomous agents**. Each class has mature tools for availability, detection, or vulnerability scanning. None answers the governing question:

| Can this asset be trusted before we act on it — and can we prove that decision?

The **GenomeX Assurance Platform** is assurance infrastructure: a vendor-agnostic overlay that evaluates, attests, and governs digital assets **before** they are restored, indexed, trained on, deployed, or executed. It sits above backup, EDR, CI/CD, and AI stacks. It does not replace them.

The architecture this product inherits is one spine, one genome, five domains:

GX-DAB → GX-SEQ → GX-SIG → GX-BE → GX-AL → domain policy
ALLOW | DENY | CONFIRM

Engine / domain	Buyer question or obligation
GX-DAB	What data exists, where it lives, which subset may proceed?
GX-SEQ	What is actually inside the asset?
GX-SIG	Does content match known-bad signatures or IOCs?
GX-BE	How has behavior drifted from the frozen baseline?
GX-AL	Can the decision be proven later?
GX-GRA	Can this recovery chain be trusted?
GX-GKA	Can this knowledge be trusted?
GX-GMA	Can this ML asset be trusted?
GX-GAA	Can this agent be trusted?
GX-GSA	Can this software release be trusted?

GX-DAB and GX-SIG are first-class platform engines. GX-GRA fuses a **Recovery Environment Genome** and a **Recovery Asset Genome** into **Restore Authority**. Every domain decision can mint a portable receipt.

Tomorrow, trust will be measured — not assumed.

1. What buyers are missing

Existing categories are strong at the questions they were built to answer. They are not pre-use assurance.

Category	Primary question	What it typically does not prove
Backup and data protection	Was data copied and retained?	Behavioral health at capture; whether restore should proceed; guest content inside the point
Storage / backup posture tools	Is the backup platform securely configured?	Whether the preserved state is clean
EDR / XDR	Is the endpoint compromised now?	Cold content at backup scale; a governed restore decision
Clean-room restore-then-scan	Is a restored copy clean?	Cost and time at nightly volume; a pre-restore gate
RAG / prompt filters	Is this query or response safe?	Whether the document should have entered the index

Category	Primary question	What it typically does not prove
AppSec / SCA	What CVEs exist?	Whether the release should be trusted given provenance and behavior

Three pains persist.

Re-infection fear. A hardened repository can faithfully preserve ransomware. A clean recovery point can sit on a drifted backup platform.


Validation does not scale. Physical clean rooms work for one critical workload. They do not rank hundreds of recovery points after an incident, or gate nightly corpora.

No portable proof. DORA, HIPAA, SEC cyber disclosure, cyber insurance, and internal audit need artifacts, not an administrator's opinion.

GenomeX is sold as the overlay that turns those investments into ALLOW / DENY / CONFIRM plus a receipt.

2. How the product is organized

GenomeX is **not** five products. It is one control plane, five shared engines, and five domain gates. Domains differ in policy objects, APIs, buyer questions, and receipts — not in duplicated worker fleets.



DAB is the librarian. The engines are the inspectors. The domains are the gates.

Decision	Meaning
ALLOW	Trust is sufficient for this use, with a receipt
DENY	Trust is insufficient; block the use
CONFIRM	Human or policy review; evidence is attached

The policy engine **escalates only**. It does not silently downgrade risk.

3. The Behavioral Genome

Buyers should not need a different trust model per domain. Every governed asset is described through eight conceptual blocks — identity, provenance, behavior, risk, drift, confidence, evidence, and attestation. Those are assurance dimensions, not eight scalar scores.

They map onto the current wire:

platform_genome_128 = 8 blocks × 16 slots

That schema is specified in the GenomeX Assurance Platform architecture paper. This product uses it on the wire. Each domain **projects** different signals into the same 128-dimensional space. A recovery chain, a document, a training set, an agent action, and a software release remain different objects. The assurance model stays consistent.

GX-BE compares the current genome to a frozen baseline, measures drift, and produces confidence. Confidence down-weights decisions when evidence is thin.

4. Shared engines

GX-DAB — Data Access Bridge

Purpose: Discover, catalog, and filter what may enter the assurance path.

What buyers get: An inventory of corpora, recovery assets, and other sources — SharePoint, object stores, filesystems, uploads — with views that emit a governed subset. DAB prevents unbounded walks from becoming the front door to restore, RAG, training, or release inspection.

Availability: **Available** for knowledge and corpus catalog; recovery-source discovery **Pilot**.

GX-SEQ — Sequencing Engine

Purpose: Open the asset and materialize truth far enough for an explainable decision.

What buyers get: Handler coverage across archives, documents, databases, VMs, containers, and artifacts; hybrid walks; backup-over-backup comparison; coverage-gated materialization; relay pull so scan fleets need not mount customer media; scan-complete evidence into fusion.

GX-SEQ does **not** own the known-bad lane. That is GX-SIG. GX-SEQ does not execute backup or restore, and it does not make the final ALLOW / DENY decision.

Availability: **Available**.

GX-SIG — Signature Engine

Purpose: Known-bad fusion — signatures, IOCs, ClamAV, YARA, and offline threat-intelligence hits applied to **materialized** bytes.

A signature hit is evidence. It is not the trust decision. Clean SIG does not prove behavioral integrity. A SIG hit on an otherwise normal genome typically raises CONFIRM or DENY under policy.

Availability: Available.

GX-BE — Behavior Engine

Purpose: Genome, baseline, drift, QSBA, and confidence at authorize time.

What buyers get: Learn → freeze → compare; Behavioral State Timeline; drift envelopes (acceptable / concerning / irreversible); explainable QSBA; domain projections of platform_genome_128 (recovery environment, recovery asset, document provenance, agent stream, software strands).

Availability: Available for GX-GRA; **Pilot** for GKA / GAA streams.

GX-AL — Assurance Ledger

Purpose: Local-first, air-gap-capable proof. Not public-blockchain marketing.

What buyers get:

Receipt	Domain	What it binds
PoS-B	GX-GRA	Environment + asset genomes, Restore Authority, Merkle root, scan outcome
KAR	GX-GKA	Document hash → extraction hash → chunk manifest → policy decision
ML-BOM	GX-GMA	Dataset hash, schema hash, training-gate decision
AAR	GX-GAA	Agent identity, action, gate decision, genome snapshot

Merkle anchoring is the integrity mechanism: a root binds the evaluated set; any later byte change invalidates the commitment. Payloads stay off-ledger.

Availability: PoS-B **Available**; KAR / AAR / ML-BOM minting **Pilot**; full ledger anchors for KAR and GSA **Roadmap**.

5. Evidence fusion

No single lane is sufficient. GenomeX fuses four lanes, plus cryptographic proof on GRA.

Lane	Question	Primary engine
Behavioral	How did the system or actor behave versus expectation?	GX-BE
Materialized content	What is inside, and was it examined deeply enough?	GX-SEQ
Anomaly	Did the asset change in a meaningful way?	GX-SEQ + GX-BE
Signature / TI	Are known indicators present?	GX-SIG
Integrity (GRA)	Can the association record be shown intact?	GX-AL Merkle + ledger

Fusion rules operators actually use:

- Worst **materialized** scan wins when several scanners report on one asset.
- Metadata-only scans inform; they do not alone authorize.
- High behavioral denial plus a SIG hit plus elevated anomaly → DENY.
- Borderline behavior with clean SIG → ALLOW or CONFIRM.
- Clean behavior with a SIG hit → CONFIRM or DENY under policy.

The Intervention Queue and Fusion Decision Panel are the operational heart of the console.

6. Assurance domains

GX-GRA — Recovery Assurance

Buyer question: *Can this recovery chain be trusted before restore?*

A recovery point can be clean while the infrastructure that created it has drifted. A hardened backup environment can faithfully preserve a compromised state. GX-GRA therefore measures both, on the same `platform_genome_128` wire, never mashed into one vector:

Recovery Environment Genome — backup application, repository, array, MFA, immutability, job-config drift, privilege, CVE posture.

Recovery Asset Genome — capture-time behavior, materialized content, signatures, provenance, proof.

Those two states fuse into **Restore Authority** — ALLOW, CONFIRM, or DENY — before any restore actuator runs. Environment drift raises ALLOW to CONFIRM; it does not, by itself, deny a clean asset. A clean environment does not authorize a compromised recovery point.

| *Trust the environment. Trust the recovery state. Then authorize the restore.*

Capability	What the buyer gets
Continuous watch	Host genome; learn, freeze, BST
Snapshot association	Vendor webhooks and local captures bound to state at time T
Dual-genome score	Environment + asset on platform_genome_128
Content assurance	GX-SEQ backup persona; materialization gate
Known-bad lane	GX-SIG on materialized guest content
Restore Authority	Tri-state authorize API
BSR	Behavioral Snapshot Recovery — snapshot graph and selective restore to a clean ancestor
Merkle integrity	Binding root and ledger anchor (not “blockchain backup”)
Time-Warp (CRIU)	Process-level checkpoint/restore, same authorize gate
Evidence export	PoS-B + ZIP pack

Do say: We prove whether restore should proceed.

Do not say: We replace Veeam, or we replace storage hardening.

Maturity: Available / Pilot.

GX-GKA — Knowledge Assurance

Buyer question: *Can this knowledge be trusted before it is indexed?*

Runtime prompt filters do not solve poisoned, sensitive, or instruction-bearing material entering the corpus. GX-DAB catalogs the estate; GX-GKA gates what may proceed to embed.

Decisions: ALLOW_INDEX, QUARANTINE, EXCLUDE_CHUNKS, REVIEW. Receipt: **KAR**.

Maturity: Pilot.

GX-GMA — Machine Learning Assurance

Buyer question: *Can this dataset or ML asset be trusted before training?*

GMA applies the same spine to JSONL, CSV, Parquet, and text training sets. Outcomes: ALLOW_TRAINING, EXCLUDE_SAMPLES, QUARANTINE_DATASET, REQUIRE_REMEDIATION. Receipt: **ML-BOM**.

GKA governs documents entering RAG. GMA governs samples entering training. Same engines; different policy objects.

Maturity: Pilot.

GX-GAA — Agent Assurance

Buyer question: *Can this agent be trusted before it acts?*

Agents push telemetry; the control plane scores and gates tool calls and delegations. GX-SEQ is not on the authorization hot path in the current design. Receipt: **AAR**.

Runtime prompt injection and tool abuse → GAA. Ingest poisoning → GKA / GMA.

Maturity: Pilot.

GX-GSA — Software Assurance

Buyer question: *Can this release be trusted before deploy?*

Five strands project into the shared genome: source control, build, dependency, release, team. SCA answers what is wrong. GSA answers whether the release should be trusted.

Maturity: Pilot.

7. Behavioral Snapshot Recovery and Merkle integrity

Behavioral Snapshot Recovery (BSR) is the GRA workbench and graph: lineage of recovery states, policy stamps (clean, suspicious, malicious), and selective restore to the nearest clean ancestor. It is a workflow, not a public chain.

Merkle anchoring is the separate integrity mechanism under GX-AL: a small on-ledger commitment (~200 bytes per snapshot), payloads off-ledger. Any later byte change is detectable by recomputation.

BSR does not claim public consensus, deduplication, off-site media replacement, or confidentiality.

GX-GRA is the product layer customers buy. BSR plus Merkle is the proof layer it depends on.

8. How GenomeX typically fits

The useful comparison is typical category capability — not a claim that every vendor product is empty on every cell.

Typical category focus	Pre-use gate	Behavioral baseline	Portable receipt	Multi-asset spine
Backup / data protection	Rarely	Repository or job telemetry	Catalog / immutability flags	Recovery
Backup posture / hardening	Config findings	Environment drift	Finding lists	Backup infrastructure
EDR rollback	Endpoint now	Host behavior	Alert / isolate	Endpoint

Typical category focus	Pre-use gate	Behavioral baseline	Portable receipt	Multi-asset spine
Clean-room scan	After restore	Limited	Report	Recovery
RAG / prompt filters	Runtime	Limited	Rarely a corpus receipt	Knowledge
AppSec / SCA	Advisory	Limited	SBOM / ticket	Software
GenomeX (this product)	Before use	platform_genome_128	PoS-B / KAR / ML-BOM / AAR	Five domains

Why buyers choose the overlay

1. Vendor-agnostic — Veeam, Cohesity, Rubrik, cloud snapshots, SharePoint, GitHub, agent frameworks remain in place.
2. Fusion, not a single scanner.
3. CONFIRM matches how SOC and DR teams work.
4. Air-gap receipts for DORA, HIPAA, insurance, and litigation.
5. One sequencing and signature investment amortized across domains.

Channel

- **MSSP:** multi-tenant overlay without owning backup media.
- **Backup OEM:** Restore Authority and PoS-B as a provable-recovery tier.
- **AI platform:** pre-index GKA before RAG goes production.
- **Insurer:** operational recoverability evidence, not a questionnaire.

Pilot success criterion: Provable ALLOW / DENY / CONFIRM on a real catalog or corpus within one business day.

9. Buyers and use cases

Buyer	Why they start here
CISO / CRO	Re-infection risk; board-ready recoverability
Backup / DR lead	Rank recovery chains without a physical clean room per point
AI platform owner	Govern corpora before embed; gate agents at action time
MSSP	Tenant posture and white-label evidence
Compliance	Exportable control operation

Buyer	Why they start here
Cyber insurer	Evidence-based recoverability and corpus risk

Domain	Representative use
GX-GRA	Ransomware ranking; Restore Authority; Time-Warp; BSR clean-ancestor restore
GX-GKA	SharePoint / S3 corpus gate; chunk poisoning; ingestion drift
GX-GMA	Training-set gate; ML-BOM for dataset attestation
GX-GAA	Tool-call gate; fleet quarantine on drift
GX-GSA	Release trust before deploy

10. What is proven and what is not

Proven today (GX-GRA Foundation): dual-genome Restore Authority; Merkle-bound snapshot association; materialized scan and GX-SIG outcomes under documented fusion; ALLOW / DENY / CONFIRM with exportable PoS-B.

Not claimed today: replacement of backup or EDR vendors; public blockchain notarization; fully automatic last-good selection without operator policy; runtime RAG filtering as a substitute for pre-index GKA; host-sensor-only continuous malware enumeration.

Air-gap: on-prem ledger, offline TI bundle, authorize path without internet.

11. Editions and maturity (September 2026)

GX-GRA edition	Capabilities
Standard	BST, QSBA, forensic queries
Advanced	Dual-genome Restore Authority, PoS-B, BSAL, hybrid TI, ledger, export (<i>Foundation today</i>)
Enterprise (roadmap)	Negative capability proofs, production ZK, precursor library, automated last-good selection

Component	Status
GX-DAB catalog and filter	Available
GX-SEQ sequencing	Available

Component	Status
GX-SIG signature / IOC lane	Available
GX-BE host / recovery genomes + QSBA	Available
GX-AL PoA ledger + Merkle + PoS-B	Available
GX-GRA Restore Authority (environment + asset)	Available / Pilot
BSR workbench (Behavioral Snapshot Recovery)	Available
Fusion UI + intervention queue	Available
GX-GKA pre-index + KAR	Pilot
GX-GMA training gate + ML-BOM	Pilot
GX-GAA action gate + AAR	Pilot
GX-GSA release trust gate	Pilot
Cross-domain posture rollout	Available
KAR / GSA full ledger anchor	Roadmap
Production ZK / PoS_N	Roadmap

Independent third-party validation is an ongoing program. Internal lab results are internal lab results.

12. Deployment

GenomeX deploys as an overlay: connectors, webhooks, and evidence ingest into the control plane; SEQ / SIG fleets adjacent to data where possible; ledger on-prem or in a sovereign enclave.

The operator console is the drill-down. Insurance, MSSP, CISO, and GRC views are read-mostly projections of the same receipts. If a claim cannot trace to a receipt field or a materialization contract, it is roadmap.

Conclusion

This paper is the product view of GenomeX. The category argument lives in *When Data Becomes Infrastructure*. The architecture this product inherits lives in *GenomeX Assurance Platform*.

What buyers get is concrete: five engines on one spine, five domain gates, dual-genome Restore Authority, four evidence lanes plus proof, and receipts — PoS-B, KAR, ML-BOM, AAR — that survive the operational moment.

Recovery should be provable, not hopeful. Across knowledge, training, software, and agents, trust should be measured before use.

| *Tomorrow, trust will be measured — not assumed.*

Knight InfoTek · GenomeX Assurance Platform · Product Whitepaper v3.0 · September 2026 ·
www.knightinfotek.com